

# Lecture 1: Theoretical foundations of ML

Pranabendu Misra  
Chennai Mathematical Institute

Advanced Machine Learning 2022

(based on slides by Madhavan Mukund)

# Supervised learning

- Set of possible input instances  $X$

# Supervised learning

- Set of possible input instances  $X$
- Categories  $C$ , say  $\{0, 1\}$

net  
 $l(n) \in C$

# Supervised learning

- Set of possible input instances  $X$
- Categories  $C$ , say  $\{0, 1\}$
- Build a classification model  $M : X \rightarrow C$

$$M(x) \rightarrow \hat{C}(x)$$

$\hat{C}(x)$

$\hat{C}(x)$

# Supervised learning

- Set of possible input instances  $X$
- Categories  $C$ , say  $\{0, 1\}$
- Build a classification model  $M : X \rightarrow C$
- Restrict the types of models
  - Hypothesis space  $\mathcal{H}$  — e.g., linear separators
  - Search for best  $M \in \mathcal{H}$

$M$  is <sup>o</sup> linear  $f^n$   
of  $n$   
VOT  $\rightarrow$  Comp.  $f^n$   
in DT

# Supervised learning

- Set of possible input instances  $X$
- Categories  $C$ , say  $\{0, 1\}$
- Build a classification model  $M : X \rightarrow C$
- Restrict the types of models
  - Hypothesis space  $\mathcal{H}$  — e.g., linear separators
  - Search for best  $M \in \mathcal{H}$
- How do we find the best  $M$ ?
  - Labelled training data (training set)
  - Choose  $M$  to minimize error (loss) with respect to the training set
  - Why should  $M$  generalize well to arbitrary data?

$$\mathcal{H} \subseteq S$$
$$l(r) \in C$$
$$S = \{ (x, l(x)) \mid x \in X \}$$


Suppose we can do this

# No free lunch

- ML algorithms minimize **training loss** (with respect to the training set)

# No free lunch

- ML algorithms minimize training loss (with respect to the training set)
- Goal is to minimize generalization loss (with respect to all inputs)




# No free lunch

- ML algorithms minimize **training** loss (with respect to the training set)
- Goal is to minimize **generalization** loss (with respect to all inputs)

## No Free Lunch Theorem [Wolpert, Macready 1997]

Averaged over all possible data distributions, every classification algorithm has the same error rate when classifying previously unobserved points.



# No free lunch

- ML algorithms minimize **training** loss (with respect to the training set)
- Goal is to minimize **generalization** loss (with respect to all inputs)

## No Free Lunch Theorem [Wolpert, Macready 1997]

Averaged over all possible data distributions, every classification algorithm has the same error rate when classifying previously unobserved points.

- Is the situation hopeless?

# No free lunch

- ML algorithms minimize **training** loss (with respect to the training set)
- Goal is to minimize **generalization** loss (with respect to all inputs)

## No Free Lunch Theorem [Wolpert, Macready 1997]

Averaged over all possible data distributions, every classification algorithm has the same error rate when classifying previously unobserved points.

- Is the situation hopeless?
- NFL theorem refers to prediction inputs coming from **all possible** distributions

# No free lunch

- ML algorithms minimize **training loss** (with respect to the training set)
- Goal is to minimize **generalization loss** (with respect to all inputs)

## No Free Lunch Theorem [Wolpert, Macready 1997]

Averaged over all possible data distributions, every classification algorithm has the same error rate when classifying previously unobserved points.

- Is the situation hopeless?
- NFL theorem refers to prediction inputs coming from **all possible** distributions
- ML assumes **training set is representative** of overall data
  - Prediction instances follow roughly the **same distribution** as training set

# A theoretical framework for ML


- $X$  is the space of input instances

# A theoretical framework for ML

- $X$  is the space of input instances
- $C \subseteq X$  is the **target concept** to be learned
  - e.g.,  $X$  is all emails,  $C$  is the set of spam emails

# A theoretical framework for ML

- $X$  is the space of input instances
- $C \subseteq X$  is the target concept to be learned
- e.g.,  $X$  is all emails,  $C$  is the set of spam emails
- $X$  is equipped with a probability distribution  $D$



$$\sum_{x \in X} D(x) = 1$$

- Any random sample from  $X$  is drawn using  $D$
- In particular, training set and test set are constituted of such random samples

$X \rightarrow$  all emails  
 $C \rightarrow$  all spam emails


# A theoretical framework for ML

- $X$  is the space of input instances
- $C \subseteq X$  is the **target concept** to be learned
  - e.g.,  $X$  is all emails,  $C$  is the set of spam emails
- $X$  is equipped with a probability distribution  $D$ 
  - Any random sample from  $X$  is drawn using  $D$
  - In particular, training set and test set are constitute of such random samples
- $\mathcal{H}$  is a set of hypotheses
  - Each  $h \in \mathcal{H}$  identifies a subset of  $X$
  - Choose the best  $h \in \mathcal{H}$  as model

# A theoretical framework for ML

- $X$  is the space of input instances
- $C \subseteq X$  is the **target concept** to be learned
  - e.g.,  $X$  is all emails,  $C$  is the set of spam emails
- $X$  is equipped with a probability distribution  $D$ 
  - Any random sample from  $X$  is drawn using  $D$
  - In particular, training set and test set are constitute of such random samples

- $\mathcal{H}$  is a set of **hypotheses**
  - Each  $h \in \mathcal{H}$  identifies a subset of  $X$
  - Choose the best  $h \in \mathcal{H}$  as model
- **True error:** Probability that  $h$  incorrectly classifies  $x \in X$  drawn randomly according to  $D$ 
  - $\text{err}_D(h) = \text{Prob}(h \Delta C)$
  - $h \Delta C = (h \setminus C) \cup (C \setminus h)$



# A theoretical framework for ML

- $X$  is the space of input instances
  - $C \subseteq X$  is the **target concept** to be learned
    - e.g.,  $X$  is all emails,  $C$  is the set of spam emails
  - $X$  is equipped with a probability distribution  $D$ 
    - Any random sample from  $X$  is drawn using  $D$
    - In particular, training set and test set are constitute of such random samples
- $\mathcal{H}$  is a set of **hypotheses**
    - Each  $h \in \mathcal{H}$  identifies a subset of  $X$
    - Choose the best  $h \in \mathcal{H}$  as model
- True error:** Probability that  $h$  incorrectly classifies  $x \in X$  drawn randomly according to  $D$
- $$\text{err}_D(h) = \text{Prob}(h \Delta C)$$
- $$h \Delta C = (h \setminus C) \cup (C \setminus h)$$
- Training error:** Given a (finite) training sample  $S \subseteq X$
- $$\text{err}_S(h) = |S \cap (h \Delta C)| / |S|$$



# A theoretical framework for ML

- $X$ , inputs with distribution  $D$
- $C \subseteq X$ , target concept
- $h \in \mathcal{H}$ , hypothesis (model) for  $C$
- True error:  $\text{err}_D(h) = \text{Prob}(h \Delta C)$
- Training error:  
 $\text{err}_S(h) = |S \cap (h \Delta C)|/|S|$

# A theoretical framework for ML

- $X$ , inputs with distribution  $D$
- $C \subseteq X$ , target concept
- $h \in \mathcal{H}$ , hypothesis (model) for  $C$
- True error:  $\text{err}_D(h) = \text{Prob}(h \Delta C)$
- Training error:  
 $\text{err}_S(h) = |S \cap (h \Delta C)|/|S|$

## Goal

Minimizing training error should correspond to minimizing true error

# A theoretical framework for ML

- $X$ , inputs with distribution  $D$
- $C \subseteq X$ , target concept
- $h \in \mathcal{H}$ , hypothesis (model) for  $C$
- True error:  $\text{err}_D(h) = \text{Prob}(h \Delta C)$
- Training error:  
 $\text{errs}(h) = |S \cap (h \Delta C)|/|S|$
- Overfitting Low training error but high true error

## Goal

Minimizing training error should correspond to minimizing true error

# A theoretical framework for ML

- $X$ , inputs with distribution  $D$
- $C \subseteq X$ , target concept
- $h \in \mathcal{H}$ , hypothesis (model) for  $C$
- True error:  $\text{err}_D(h) = \text{Prob}(h \Delta C)$
- Training error:  
 $\text{errs}(h) = |S \cap (h \Delta C)|/|S|$
- Overfitting Low training error but high true error
- Underfitting Cannot achieve low training/true error

## Goal

Minimizing training error should correspond to minimizing true error

# A theoretical framework for ML

- $X$ , inputs with distribution  $D$
- $C \subseteq X$ , target concept
- $h \in \mathcal{H}$ , hypothesis (model) for  $C$
- True error:  $\text{err}_D(h) = \text{Prob}(h \Delta C)$
- Training error:  
 $\text{errs}(h) = |S \cap (h \Delta C)|/|S|$

## Goal

Minimizing training error should correspond to minimizing true error

- **Overfitting** Low training error but high true error
- **Underfitting** Cannot achieve low training/true error
- Related to the representational capacity of  $\mathcal{H}$ 
  - How expressive is  $\mathcal{H}$ ? How many different concepts can it capture?
  - Capacity too high — overfitting
  - Capacity too low — underfitting

# Probably Approximately Correct (PAC) learning

- Assume  $\mathcal{H}$  is finite — use  $|\mathcal{H}|$  for capacity

# Probably Approximately Correct (PAC) learning

- Assume  $\mathcal{H}$  is finite — use  $|\mathcal{H}|$  for capacity

- Probably Approximately Correct learning

With high probability, the hypothesis  $h$  that fits the sample  $S$  also fits the concept  $C$  approximately correctly

burst  
ht  $H$   
wnt  $S$   
learnt trainings

# Probably Approximately Correct (PAC) learning

- Assume  $\mathcal{H}$  is finite — use  $|\mathcal{H}|$  for capacity
- Probably Approximately Correct learning

With high probability, the hypothesis  $h$  that fits the sample  $S$  also fits the concept  $C$  approximately correctly

## Theorem (PAC learning guarantee)

Let  $\delta, \epsilon > 0$ . Let  $S$  be a training set of size  $n \geq \frac{1}{\epsilon}(\ln |\mathcal{H}| + \ln(1/\delta))$  drawn using  $D$ .  
With probability  $\geq 1 - \delta$ , every  $h \in \mathcal{H}$  with training error zero has true error  $< \epsilon$

- Size of the sample required for PAC guarantee determined by parameters  $\delta, \epsilon$ 
  - Smaller  $\delta$  means higher probability of find a good hypothesis
  - Smaller  $\epsilon$  means better performance with respect to generalization



# Probably Approximately Correct (PAC) learning

## Theorem (Uniform convergence)

Let  $\delta, \epsilon > 0$ . Let  $S$  be a training set of size  $n \geq \frac{1}{2\epsilon^2}(\ln |\mathcal{H}| + \ln(2/\delta))$  drawn using  $D$ . With probability  $\geq 1 - \delta$ , every  $h \in \mathcal{H}$  satisfies  $|\text{err}_S(h) - \text{err}_D(h)| \leq \epsilon$ .

- Stronger guarantee: even if we cannot achieve zero training error, the additional generalization error is bounded

# Probably Approximately Correct (PAC) learning

## Theorem (Uniform convergence)

Let  $\delta, \epsilon > 0$ . Let  $S$  be a training set of size  $n \geq \frac{1}{2\epsilon^2}(\ln |\mathcal{H}| + \ln(2/\delta))$  drawn using  $D$ . With probability  $\geq 1 - \delta$ , every  $h \in \mathcal{H}$  satisfies  $|\text{err}_S(h) - \text{err}_D(h)| \leq \epsilon$ .

- Stronger guarantee: even if we cannot achieve zero training error, the additional generalization error is bounded
- What if  $\mathcal{H}$  is not finite?

# Probably Approximately Correct (PAC) learning

## Theorem (Uniform convergence)

Let  $\delta, \epsilon > 0$ . Let  $S$  be a training set of size  $n \geq \frac{1}{2\epsilon^2}(\ln |\mathcal{H}| + \ln(2/\delta))$  drawn using  $D$ . With probability  $\geq 1 - \delta$ , every  $h \in \mathcal{H}$  satisfies  $|\text{err}_S(h) - \text{err}_D(h)| \leq \epsilon$ .

↑ VC-Dim( $\mathcal{H}$ )

- Stronger guarantee: even if we cannot achieve zero training error, the additional generalization error is bounded
- What if  $\mathcal{H}$  is not finite?
- Other measures of capacity — e.g. VC-dimension

of  $\mathcal{H}$

# Probably Approximately Correct (PAC) learning

## Theorem (Uniform convergence)

Let  $\delta, \epsilon > 0$ . Let  $S$  be a training set of size  $n \geq \frac{1}{2\epsilon^2}(\ln |\mathcal{H}| + \ln(2/\delta))$  drawn using  $D$ . With probability  $\geq 1 - \delta$ , every  $h \in \mathcal{H}$  satisfies  $|\text{err}_S(h) - \text{err}_D(h)| \leq \epsilon$ .

- Stronger guarantee: even if we cannot achieve zero training error, the additional generalization error is bounded
- What if  $\mathcal{H}$  is not finite?
- Other measures of capacity — e.g. VC-dimension
- Analogous convergence theorems in terms of VC-dimension

# Overfitting and underfitting

## Example: Regression

- $\mathcal{H}_d$  is set of polynomials of degree  $d$

# Overfitting and underfitting

## Example: Regression

- $\mathcal{H}_d$  is set of polynomials of degree  $d$
- Increasing  $d$  increases expressiveness — higher representational capacity

# Overfitting and underfitting

## Example: Regression

- $\mathcal{H}_d$  is set of polynomials of degree  $d$
- Increasing  $d$  increases expressiveness — higher representational capacity
- Using too high a  $d$  results in overfitting

# Overfitting and underfitting


## Example: Regression

- $\mathcal{H}_d$  is set of polynomials of degree  $d$
- Increasing  $d$  increases expressiveness — higher representational capacity
- Using too high a  $d$  results in overfitting
- Using too low a  $d$  results in underfitting

# Overfitting and underfitting

## Example: Regression


- $\mathcal{H}_d$  is set of polynomials of degree  $d$
- Increasing  $d$  increases expressiveness — higher representational capacity
- Using too high a  $d$  results in overfitting
- Using too low a  $d$  results in underfitting



# Overfitting and underfitting

## Example: Regression

- $\mathcal{H}_d$  is set of polynomials of degree  $d$
- Increasing  $d$  increases expressiveness — higher representational capacity
- Using too high a  $d$  results in overfitting
- Using too low a  $d$  results in underfitting




- Random points lying along a quadratic
- Linear function underfits

# Overfitting and underfitting

## Example: Regression

- $\mathcal{H}_d$  is set of polynomials of degree  $d$
- Increasing  $d$  increases expressiveness — higher representational capacity
- Using too high a  $d$  results in overfitting
- Using too low a  $d$  results in underfitting

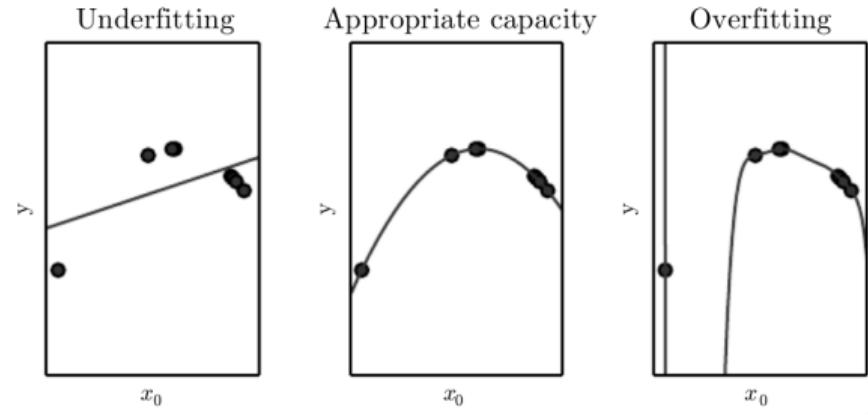


- Random points lying along a quadratic
- Linear function underfits
- Quadratic fits and generalizes well

# Overfitting and underfitting



## Example: Regression

- $\mathcal{H}_d$  is set of polynomials of degree  $d$
- Increasing  $d$  increases expressiveness — higher representational capacity
- Using too high a  $d$  results in overfitting
- Using too low a  $d$  results in underfitting




- Random points lying along a quadratic
- Linear function underfits
- Quadratic fits and generalizes well
- Degree 9 polynomial overfits

# Capacity and error




- As capacity increases, training error decreases

# Capacity and error




- As capacity increases, training error decreases
- Initially, generalization error also decreases

# Capacity and error



- As capacity increases, training error decreases
- Initially, generalization error also decreases
- At some point, generalization error starts increasing

# Capacity and error



- As capacity increases, training error decreases
- Initially, generalization error also decreases
- At some point, generalization error starts increasing
- Optimum capacity is not where training error is minimum

# Theory and practice

- Deep learning models are too complex to compute representational capacity explicitly

# Theory and practice

- Deep learning models are too complex to compute representational capacity explicitly
- May not even be able to achieve true representational capacity

# Theory and practice

- Deep learning models are too complex to compute representational capacity explicitly
- May not even be able to achieve true representational capacity
- Effective capacity limited by capabilities of parameter estimation algorithm (backpropagation with optimization)

# Theory and practice

- Deep learning models are too complex to compute representational capacity explicitly
- May not even be able to achieve true representational capacity
- Effective capacity limited by capabilities of parameter estimation algorithm (backpropagation with optimization)
- Parameter estimation is a complex nonlinear optimization

# Theory and practice

- Deep learning models are too complex to compute representational capacity explicitly
- May not even be able to achieve true representational capacity
- Effective capacity limited by capabilities of parameter estimation algorithm (backpropagation with optimization)
- Parameter estimation is a complex nonlinear optimization

## Regularization

- Add a penalty for model complexity to the loss function
- Trade off lower training error against penalty

$$\text{Loss}(h) = \frac{\sum \text{[if } \text{out} \text{ is } \text{wrong} \text{ or } \text{r} \text{]}}{|\text{S}|} + \text{Rg Tm } e^{(\text{off} + \text{out})}$$

$\gamma = mn + c$   
↑ how I choose m or ↓ c

# Theory and practice

- Deep learning models are too complex to compute representational capacity explicitly
- May not even be able to achieve true representational capacity
- Effective capacity limited by capabilities of parameter estimation algorithm (backpropagation with optimization)
- Parameter estimation is a complex nonlinear optimization

## Regularization

- Add a penalty for model complexity to the loss function
- Trade off lower training error against penalty

## Hyperparameters

- Settings that adjust the capacity — e.g., degree of polynomial
- Set externally, not learned
- Search hyperparameter combinations for optimal settings

# Summary

- Supervised learning builds a model that minimize training error

# Summary

- Supervised learning builds a model that minimize training error
- Real goal is to minimize generalization error

# Summary

- Supervised learning builds a model that minimize training error
- Real goal is to minimize generalization error
- PAC learning provides a theoretical framework to justify this

- Supervised learning builds a model that minimize training error
- Real goal is to minimize generalization error
- PAC learning provides a theoretical framework to justify this
- Discrepancies in representational capacity of models can cause underfitting or overfitting

# Summary

- Supervised learning builds a model that minimize training error
- Real goal is to minimize generalization error
- PAC learning provides a theoretical framework to justify this
- Discrepancies in representational capacity of models can cause underfitting or overfitting
- In practice, use regularization and hyperparameter search to identify optimum capacity

Rec from  
help w/  
w/